

---

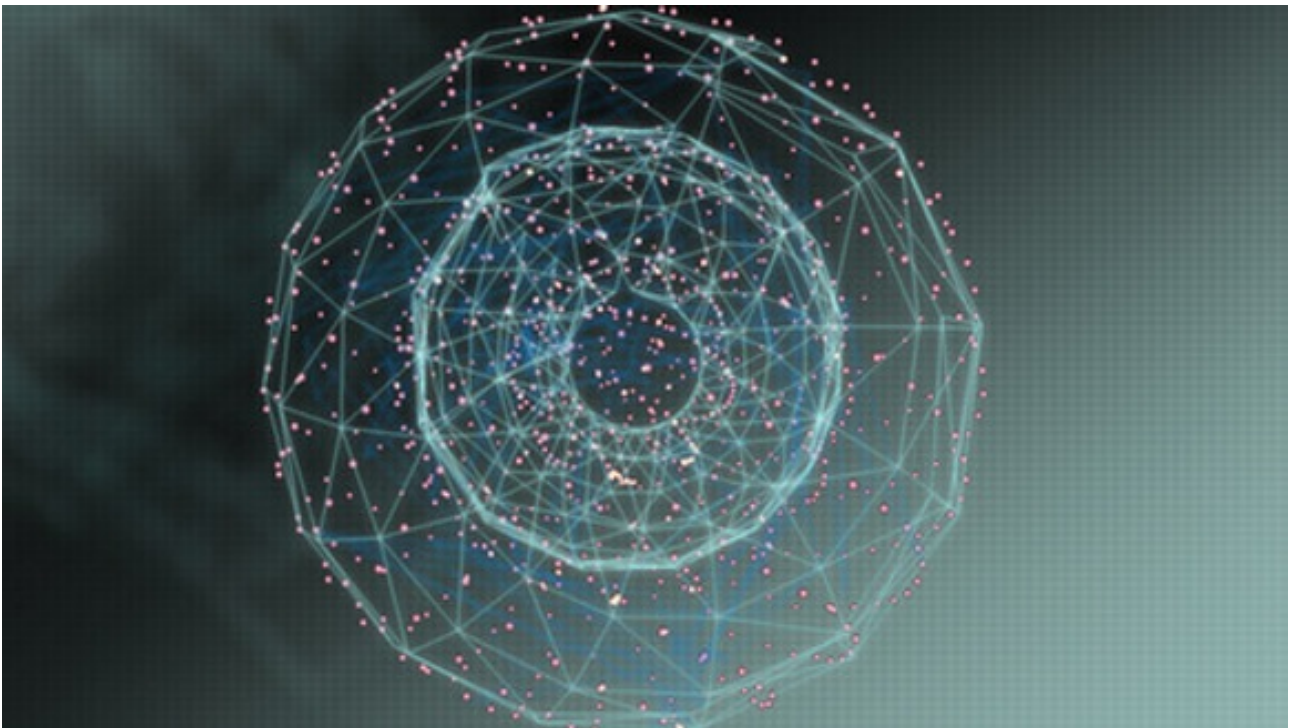
# PATTERN DISCOVERY IN DATA MINING

---

Concepts and challenges in pattern discovery and analysis.  
Pattern evaluation, mining and classification

Course author:

JIAWEI HAN



*University of Illinois at Urbana-Champaign  
&  
Coursera*

2015

# Contents

<b>1</b>	<b>Lecture 2: Pattern Discovery Basic Concepts</b>	<b>3</b>
1.1	Frequent Itemsets (Patterns)	3
1.2	Association Rules	3
1.3	Expressing Patterns in Compressed Form	3
1.4	Recommended readings	3
<b>2</b>	<b>Lecture 3. Efficient Pattern Mining Methods</b>	<b>4</b>
2.1	The Downward Closure Property of Frequent Patterns	4
2.2	The Apriori Algorithm	4
2.2.1	Algorithm pseudocode	4
2.2.2	How to generate candidates?	4
2.3	Extensions or Improvements of Apriori	5
2.3.1	Partitioning	5
2.3.2	Direct Hashing and Pruning (DHP)	5
2.4	Vertical Data Format	5
2.5	A Pattern Growth Approach	6
2.6	CLOSET+: Mining Closed Itemsets by Pattern-Growth	6
2.7	Recommended readings	6
<b>3</b>	<b>Lecture 4: Pattern Evaluation</b>	<b>7</b>
3.1	Interestingness Measures: Lift and $\chi^2$	7
3.1.1	Interestingness Measure: Lift	7
3.1.2	Interestingness Measure: $\chi^2$	7
3.2	Null Invariance Measures	8
3.3	Imbalance Ratio	8
3.4	Recommended Readings	8
<b>4</b>	<b>Lecture 4: Mining Diverse Patterns</b>	<b>8</b>
4.1	Mining Multi-Level Associations	8
4.2	Mining Multi-Dimensional Associations	9
4.3	Mining Quantitative Associations	9
4.4	Mining Negative Correlations	9
4.5	Mining Compressed Patterns	10
4.5.1	Mining Compressed Patterns	10
4.5.2	Redundancy-Aware Top-k Patterns	10
4.6	Mining Colossal Patterns	10
4.6.1	Pattern-Fusion	10
4.6.2	Robustness of Colossal Patterns	11
4.6.3	The Pattern-Fusion Algorithm	11
4.7	Recommended Readings	12
<b>5</b>	<b>Constraint-Based Pattern Mining</b>	<b>12</b>
5.1	Meta-Rule Guided Mining	12
5.2	Kinds of Constraints	12
5.2.1	Pattern space pruning constraints	13
5.2.2	Data space pruning constraints	13
5.3	Recommended Readings	13

<b>6</b>	<b>Sequential Pattern Mining</b>	<b>13</b>
6.1	Sequential Pattern . . . . .	13
6.2	GSP: Apriori-Based Sequential Pattern Mining . . . . .	14
6.3	SPADE: Sequential Pattern Mining in Vertical Data Format . . . . .	14
6.4	PrefixSpan: A Pattern-Growth Approach . . . . .	15
6.5	CloSpan: Mining Closed Sequential Patterns . . . . .	15
6.6	Constraint-Based Sequential-Pattern Mining . . . . .	16
6.6.1	Timing-Based Constraints . . . . .	16
6.7	Recommended Readings . . . . .	16
<b>7</b>	<b>Lecture 8. Graph Pattern Mining</b>	<b>17</b>
7.1	Frequent (Sub)Graph Patterns . . . . .	17
7.2	Apriori-Based Approach . . . . .	18
7.3	gSPAN: Graph Pattern Growth . . . . .	18
7.4	Mining Closed Graph Patterns . . . . .	18
7.5	gIndex: A Graph Indexing Method . . . . .	19
7.6	SpiderMine: Mining Top-K Large Structural Patterns in a Massive Network .	19
7.7	Recommended Readings . . . . .	20
<b>8</b>	<b>Lecture 9. Pattern-Based Classification</b>	<b>21</b>
8.1	Classification: Basic Concepts . . . . .	21
8.2	Pattern-based classification methods . . . . .	21
8.3	CBA: Classification Based on Associations . . . . .	22
8.4	CMAR: Classification Based on Multiple Association Rules . . . . .	22
8.5	Discriminative Pattern-Based Classification . . . . .	22
8.5.1	On the Power of Discriminative Patterns . . . . .	23
8.6	DDPMine: Direct Mining of Discriminative Patterns . . . . .	23
8.6.1	Branch-and-Bound Search . . . . .	23
8.7	Recommended Readings . . . . .	24
<b>9</b>	<b>Lecture 10. Exploring Pattern Mining Applications</b>	<b>24</b>
9.1	Strategy 1: Simultaneously Inferring Phrases and Topics . . . . .	25
9.2	Strategy 2: Post Topic Modeling Phrase Construction . . . . .	25
9.2.1	TurboTopics . . . . .	25
9.2.2	KERT . . . . .	25
9.3	Strategy 3: First Phrase Mining then Topic Modeling . . . . .	26
9.3.1	Phrase Mining: Frequent Pattern Mining + Statistical Analysis . . . . .	26
9.4	Recommended Readings . . . . .	27
<b>10</b>	<b>Lecture 11: Advanced Topics on Pattern Discovery</b>	<b>27</b>
10.1	Frequent Pattern Mining in Data Streams . . . . .	27
10.1.1	Lossy Counting Algorithm . . . . .	27
10.1.2	Recommended Readings . . . . .	28
10.2	Spatiotemporal and Trajectory Pattern Mining . . . . .	28
10.2.1	Mining Relative Movement Patterns . . . . .	28
10.2.2	Recommended Readings . . . . .	28
10.3	Pattern Discovery for Software Bug Mining . . . . .	29
10.3.1	Typical Software Bug Detection Methods . . . . .	29
10.3.2	Mining Copy-and-Paste Bugs . . . . .	29
10.4	Pattern Discovery for Image Analysis . . . . .	30
10.4.1	Recommended Readings . . . . .	30

# 1 Lecture 2: Pattern Discovery Basic Concepts

## 1.1 Frequent Itemsets (Patterns)

$X$  = itemset

- **(absolute) support (count) of  $X$ :** Frequency or the number of occurrences of an itemset  $X$
- **(relative) support,  $s$ :** The fraction of transactions that contains  $X$  (i.e., the probability that a transaction contains  $X$ )
- An itemset  $X$  is **frequent** if the support of  $X$  is no less than a *minsup* threshold (denoted as  $\sigma$ ):  $\text{sup}(X) \geq \sigma$ .

## 1.2 Association Rules

Association rules:  $X \rightarrow Y(s, c)$ :

- **Support ( $s$ ):** the probability that a transaction contains  $X \cup Y$ :

$$\text{sup}(X \rightarrow Y) = P(X \cup Y)$$

- **Confidence ( $c$ ):** the conditional probability that a transaction containing  $X$  also contains  $Y$ :

$$c = P(Y | X) = \frac{\text{sup}(X \cup Y)}{\text{sup}(X)}$$

## 1.3 Expressing Patterns in Compressed Form

**Definition. Closed patterns:** A pattern (itemset)  $X$  is closed if  $X$  is frequent, and there exists no super-pattern  $Y \supset X$ , with the same support as  $X$ .

Closed pattern is a lossless compression of frequent patterns.

**Definition. Max-patterns:** A pattern  $X$  is a max-pattern if  $X$  is frequent and there exists no frequent super-pattern  $Y \supset X$ .

Max-pattern is a lossy compression!

## 1.4 Recommended readings

- R. Agrawal, T. Imielinski, and A. Swami, «Mining association rules between sets of items in large databases», in Proc. of SIGMOD'93
- R. J. Bayardo, «Efficiently mining long patterns from databases», in Proc. of SIGMOD'98
- N. Pasquier, Y. Bastide, R. Taouil, and L. Lakhal, «Discovering frequent closed itemsets for association rules», in Proc. of ICDT'99
- J. Han, H. Cheng, D. Xin, and X. Yan, «Frequent Pattern Mining: Current Status and Future Directions», Data Mining and Knowledge Discovery, 15(1): 55-86, 2007

## 2 Lecture 3. Efficient Pattern Mining Methods

### 2.1 The Downward Closure Property of Frequent Patterns

The downward closure (also called «Apriori») property of frequent patterns: **Any subset of a frequent itemset must be frequent.** Apriori pruning principle: **If there is any itemset which is infrequent, its superset should not even be generated!**

Scalable mining Methods: Three major approaches

- Level-wise, join-based approach: Apriori (2.2)
- Vertical data format approach: Eclat (2.4)
- Frequent pattern projection and growth: FPgrowth (2.5)

### 2.2 The Apriori Algorithm

#### 2.2.1 Algorithm pseudocode

$C_k$ : Candidate itemset of size  $k$

$F_k$ : Frequent itemset of size  $k$

TDB = transactional database

---

#### Algorithm 1 The Apriori Algorithm

---

```
 $k := 1$   
 $F_k :=$  frequent items # frequent 1-itemset  
while  $F_k \neq \emptyset$  do  
     $C_{k+1} :=$  candidates generated from  $F_k$  # candidate generation  
    Derives  $F_{k+1}$  by counting candidates in  $C_{k+1}$  with respect to TDB at  
    minsup  
     $k := k + 1$   
end while  
return  $\cup_k F_k$  # return  $F_k$  generated at each level
```

---

#### 2.2.2 How to generate candidates?

- Step1: self-joining  $F_k$
- Step2: pruning

---

#### Algorithm 2 Step1: self-joining $F_k$

---

```
insert into  $C_k$   
select  $p.item_1, p.item_2, \dots, p.item_{k-1}, q.item_{k-1}$   
from  $F_{k-1}$  as  $p, F_{k-1}$  as  $q$   
where  $p.item_1 = q.item_1, \dots, p.item_{k-2} = q.item_{k-2}, p.item_{k-1} < q.item_{k-1}$ 
```

---

---

**Algorithm 3** Step2: pruning

---

```
for all itemsets  $c$  in  $C_k$  do
  for all  $(k-1)$  subsets  $s$  of  $c$  do
    if  $s$  is not in  $F_{k-1}$  then
      delete  $c$  from  $C_k$ 
    end if
  end for
end for
```

---

## 2.3 Extensions or Improvements of Apriori

- Reduce passes of transaction database scans
  - Partitioning
  - Dynamic itemset counting
- Shrink the number of candidates
  - Hashing
  - Pruning by support lower bounding
  - Sampling
- Exploring special data structures
  - Tree projection
  - H-miner
  - Hypercube decomposition

### 2.3.1 Partitioning

**Theorem.** Any itemset that is potentially frequent in TDB must be frequent in at least one of the partitions of TDB

Method: Scan Database Only Twice:

- Scan 1: Partition database (how?) and find local frequent patterns
- Scan 2: Consolidate global frequent patterns (how to?)

### 2.3.2 Direct Hashing and Pruning (DHP)

Observation: A  $k$ -itemset whose corresponding hashing bucket count is below the threshold cannot be frequent

## 2.4 Vertical Data Format

**ECLAT** - Equivalence Class Transformation

Frequent patterns are derived based on vertical intersections. To accelerate data mining you can use **diffset**: only keep track of differences of tids.

## 2.5 A Pattern Growth Approach

FP-tree - frequent pattern tree

TID	Items in the Transaction	Ordered, frequent items
100	{f, a, c, d, g, i, m, p}	{f, c, a, m, p}
200	{a, b, c, f, l, m, o}	{f, c, a, b, m}
300	{b, f, h, j, o, w}	{f, b}
400	{b, c, k, s, p}	{c, b, p}
500	{a, f, c, e, l, p, m, n}	{f, c, a, m, p}

Figure 1: Transactional DB

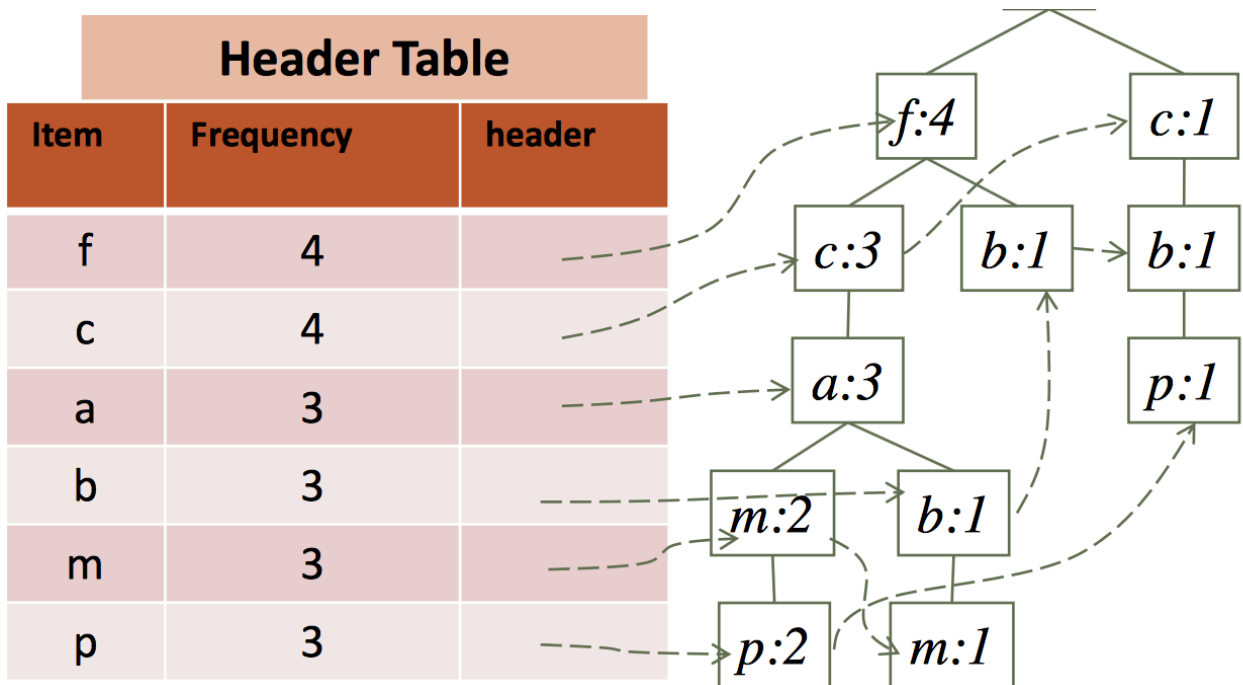


Figure 2: FP-tree

## 2.6 CLOSET+: Mining Closed Itemsets by Pattern-Growth

Itemset merging: *If Y appears in every occurrence of X, then Y is merged with X*

## 2.7 Recommended readings

- R. Agrawal and R. Srikant, «Fast algorithms for mining association rules», VLDB'94
- A. Savasere, E. Omiecinski, and S. Navathe, «An efficient algorithm for mining association rules in large databases», VLDB'95
- J. S. Park, M. S. Chen, and P. S. Yu, «An effective hash-based algorithm for mining association rules», SIGMOD'95

- S. Sarawagi, S. Thomas, and R. Agrawal, «Integrating association rule mining with relational database systems: Alternatives and implications», SIGMOD'98
- M. J. Zaki, S. Parthasarathy, M. Ogihara, and W. Li, «Parallel algorithm for discovery of association rules», Data Mining and Knowledge Discovery, 1997
- J. Han, J. Pei, and Y. Yin, «Mining frequent patterns without candidate generation», SIGMOD'00
- M. J. Zaki and Hsiao, «CHARM: An Efficient Algorithm for Closed Itemset Mining», SDM'02
- J. Wang, J. Han, and J. Pei, «CLOSET+: Searching for the Best Strategies for Mining Frequent Closed Itemsets», KDD'03
- C. C. Aggarwal, M.A., Bhuiyan, M. A. Hasan, «Frequent Pattern Mining Algorithms: A Survey», in Aggarwal and Han (eds.): Frequent Pattern Mining, Springer, 2014

### 3 Lecture 4: Pattern Evaluation

#### 3.1 Interestingness Measures: Lift and $\chi^2$

##### 3.1.1 Interestingness Measure: Lift

**Lift** - measure of dependent/correlated events:

$$\text{lift}(B, C) = \frac{c(B \rightarrow C)}{s(C)} = \frac{s(B \cup C)}{s(B) \times s(C)}$$

Lift(B, C) may tell how B and C are correlated:

- $\text{Lift}(B, C) = 1$ : B and C are independent
- $\text{Lift}(B, C) > 1$ : positively correlated
- $\text{Lift}(B, C) < 1$ : negatively correlated

##### 3.1.2 Interestingness Measure: $\chi^2$

$$\chi^2 = \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}}$$

General rules:

- $\chi^2 = 0$ : independent
- $\chi^2 > 0$ : correlated, either positive or negative, so it needs additional test

Too many null transactions may lead to invalid correlation result!



## 3.2 Null Invariance Measures

$$\begin{aligned}\text{AllConf}(A, B) &= \frac{s(A \cup B)}{\max\{s(A), s(B)\}} \\ \text{Jaccard}(A, B) &= \frac{s(A \cup B)}{s(A) + s(B) - s(A \cup B)} \\ \text{Cosine}(A, B) &= \frac{s(A \cup B)}{\sqrt{s(A) \times s(B)}} \\ \text{Kulczynsky}(A, B) &= \frac{1}{2} \left( \frac{s(A \cup B)}{s(A)} + \frac{s(A \cup B)}{s(B)} \right) \\ \text{MacConf}(A, B) &= \max \left\{ \frac{s(A)}{s(A \cup B)}, \frac{s(B)}{s(A \cup B)} \right\}\end{aligned}$$

## 3.3 Imbalance Ratio

IR (Imbalance Ratio): measure the imbalance of two itemsets A and B in rule implications:

$$\text{IR}(A, B) = \frac{|s(A) - s(B)|}{s(A) + s(B) - s(A \cup B)}$$

Kulczynski and Imbalance Ratio (IR) together present a clear picture

## 3.4 Recommended Readings

- C. C. Aggarwal and P. S. Yu. A New Framework for Itemset Generation. PODS'98
- S. Brin, R. Motwani, and C. Silverstein. Beyond market basket: Generalizing association rules to correlations. SIGMOD'97
- M. Klemettinen, H. Mannila, P. Ronkainen, H. Toivonen, and A. I. Verkamo. Finding interesting rules from large sets of discovered association rules. CIKM'94
- E. Omiecinski. Alternative Interest Measures for Mining Associations. TKDE'03
- P.-N. Tan, V. Kumar, and J. Srivastava. Selecting the Right Interestingness Measure for Association Patterns. KDD'02
- T. Wu, Y. Chen and J. Han, Re-Examination of Interestingness Measures in Pattern Mining: A Unified Framework, Data Mining and Knowledge Discovery, 21(3):371-397, 2010

# 4 Lecture 4: Mining Diverse Patterns

## 4.1 Mining Multi-Level Associations

Items often form hierarchies. How to set min-support thresholds? **Level-reduced min-support:** items at the lower level are expected to have lower support.

Efficient mining: **shared** multi-level mining. Use the lowest min-support to pass down the set of candidates.

Redundancy<sup>1</sup> filtering: some rules may be redundant due to «ancestor»<sup>2</sup> relationships between items. A rule is **redundant** if:

- its support is close to the «expected» value, according to its «ancestor» rule
- it has a similar confidence as its «ancestor».

It is necessary to have customized min-support settings for different kinds of items: group-based «individualized» min-support.

## 4.2 Mining Multi-Dimensional Associations

Rules can be single-dimensional or multi-dimensional:

- Single-dimensional:

$$\text{buys}(X, \text{«milk»}) \Rightarrow \text{buys}(X, \text{«bread»})$$

- Inter-dimension association rule:

$$\text{age}(X, \text{«18-25»}) \wedge \text{occupation}(X, \text{«student»}) \Rightarrow \text{buys}(X, \text{«coke»})$$

- Hybrid-dimension association rules:

$$\text{age}(X, \text{«18-25»}) \wedge \text{buys}(X, \text{«popcorn»}) \Rightarrow \text{buys}(X, \text{«coke»})$$

Attributes can be categorical or numerical

## 4.3 Mining Quantitative Associations

Methods:

- Static discretization based on predefined concept hierarchies
- Dynamic discretization based on data distribution
- Clustering: distance-based association
- Deviation analysis

## 4.4 Mining Negative Correlations

- Rare patterns = very low support but interesting
- Negative patterns = negatively correlated, unlikely to happen together

A support-based definition: if itemsets A and B are both frequent but rarely occur together, i.e.,  $\text{sup}(A \cup B) \ll \text{sup}(A) \times \text{sup}(B)$  then A and B are negatively correlated.

The support-based definition is not null-invariant!

A Kulczynski measure-based definition: if itemsets A and B are frequent but  $\frac{P(A|B) + P(B|A)}{2} < \varepsilon$ , where  $\varepsilon$  is a negative pattern threshold, then A and B are negatively correlated.

---

<sup>1</sup>Redundancy - избыточность

<sup>2</sup>Ancestor – предок

## 4.5 Mining Compressed Patterns

### 4.5.1 Mining Compressed Patterns

Pattern distance measure:

$$Dist(P_1, P_2) = 1 - \frac{|T(P_1) \cap T(P_2)|}{|T(P_1) \cup T(P_2)|}$$

**$\delta$ -clustering.** For each pattern P, find all patterns which can be expressed by P and whose distance to P is within  $\delta$  ( $\delta$ -cover). All patterns in the cluster can be represented by P = compressed patterns.<sup>3</sup>

### 4.5.2 Redundancy-Aware Top-k Patterns

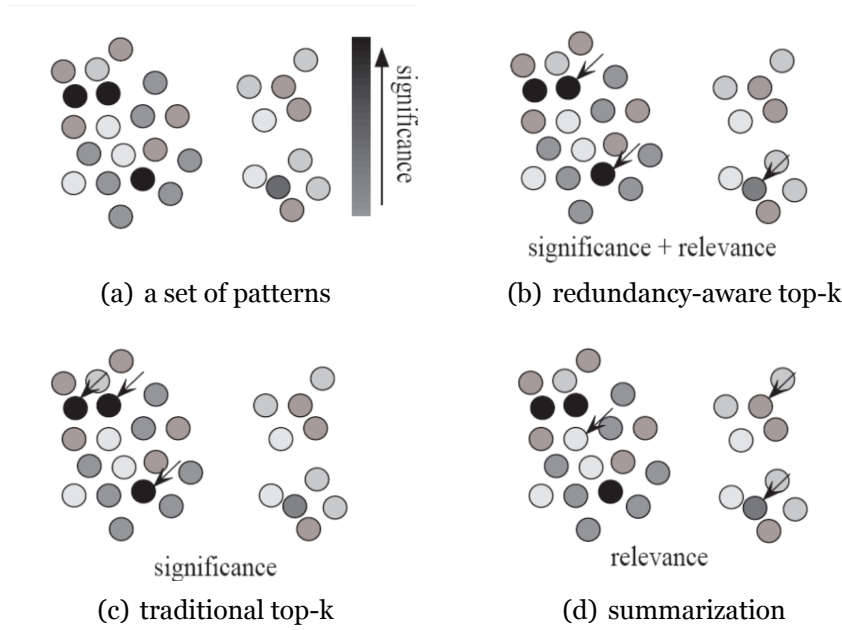


Figure 3: Desired patterns: high significance & low redundancy

Use **MMS (Maximal Marginal Significance)** for measuring the combined significance of a pattern set.<sup>4</sup>

## 4.6 Mining Colossal Patterns

### 4.6.1 Pattern-Fusion

**Pattern fusion strategy:** fuse small patterns together in one step to generate new pattern candidates of significant sizes.

Subpatterns  $\alpha_1$  to  $\alpha_k$  cluster tightly around the colossal pattern  $\alpha$  by sharing a similar support. Such subpatterns are **core patterns** of  $\alpha$ . A colossal pattern can be generated by merging a set of core patterns.

<sup>3</sup>Method for efficient, direct mining of compressed frequent patterns: Xin et al., VLDB'05.

<sup>4</sup>Xin et al., Extracting Redundancy-Aware Top-K Patterns, KDD'06.

#### 4.6.2 Robustness of Colossal Patterns

**Definition.** For a frequent pattern  $\alpha$ , a subpattern  $\beta$  is a  $\tau$ -core pattern of  $\alpha$  if  $\beta$  shares a similar support set with  $\alpha$ , i.e.,

$$\frac{|D_\alpha|}{|D_\beta|} \geq \tau, 0 < \tau \leq 1,$$

where  $\tau$  is called the **core ratio**.

**Definition.**  $(d, \tau)$ -robustness<sup>5</sup>: a pattern  $\alpha$  is  $(d, \tau)$ -robust if  $d$  is the maximum number of items that can be removed from  $\alpha$  for the resulting pattern to remain a  $\tau$ -core pattern of  $\alpha$ :

$$d = \max_{\beta} \{|\alpha| - |\beta| \mid \beta \subseteq \alpha, \text{ and } \beta \text{ is a } \tau\text{-core pattern of } \alpha\}$$

For a pattern  $\alpha$  let  $C_\alpha$  be the set of all its core patterns for a specified  $\tau$ :

$$C_\alpha = \{\beta \mid \beta \subseteq \alpha, \frac{|D_\alpha|}{|D_\beta|} \geq \tau\}$$

**Theorem.** For a  $(d, \tau)$ -robust pattern  $\alpha$ :

$$|C_\alpha| \geq 2^d$$

**Robustness of Colossal Patterns:** a colossal pattern tends to have much more core patterns than small patterns. Such core patterns can be clustered together to form «dense balls» based on pattern distance defined by

$$Dist(\alpha, \beta) = 1 - \frac{|D_\alpha \cap D_\beta|}{|D_\alpha \cup D_\beta|}$$

**Theorem.** For two patterns  $\beta_1, \beta_2 \in C_\alpha$

$$Dist(\beta_1, \beta_2) \leq r(\tau), \text{ where } r(\tau) = 1 - \frac{1}{2/\tau - 1}$$

#### 4.6.3 The Pattern-Fusion Algorithm

- Initialization (Creating initial pool): Use an existing algorithm to mine all frequent patterns up to a small size, e.g., 3
- Iteration (Iterative Pattern Fusion):
  - At each iteration,  $K$  seed patterns are randomly picked from the current pattern pool
  - For each seed pattern thus picked, we find all the patterns within a bounding ball centered at the seed pattern
  - All these patterns found are fused together to generate a set of super-patterns
  - All the super-patterns thus generated form a new pool for the next iteration
- Termination: when the current pool contains no more than  $K$  patterns at the beginning of an iteration

---

<sup>5</sup>Robustness - прочность

## 4.7 Recommended Readings

- R. Srikant and R. Agrawal, «Mining generalized association rules», VLDB'95
- Y. Aumann and Y. Lindell, «A Statistical Theory for Quantitative Association Rules», KDD'99
- D. Xin, J. Han, X. Yan and H. Cheng, «On Compressing Frequent Patterns», Knowledge and Data Engineering, 60(1): 5-29, 2007
- D. Xin, H. Cheng, X. Yan, and J. Han, «Extracting Redundancy-Aware Top-K Patterns», KDD'o6
- F. Zhu, X. Yan, J. Han, P. S. Yu, and H. Cheng, «Mining Colossal Frequent Patterns by Core Pattern Fusion», ICDE'o7
- J. Han, H. Cheng, D. Xin, and X. Yan, «Frequent Pattern Mining: Current Status and Future Directions», Data Mining and Knowledge Discovery, 15(1): 55-86, 2007

## 5 Constraint-Based Pattern Mining

### 5.1 Meta-Rule Guided Mining

In general, (meta) rules can be in the form of

$$P_1 \wedge P_2 \wedge \dots \wedge P_l \Rightarrow Q_1 \wedge Q_2 \wedge \dots \wedge Q_r$$

Method to find meta-rules:

- Find frequent ( $l + r$ ) predicates (based on min-support)
- Push constraints deeply when possible into the mining process
- Also, push min\_conf, min\_correlation, and other measures as early as possible (measures acting as constraints)

### 5.2 Kinds of Constraints

- Pattern space pruning constraints
  - Anti-monotonic: If constraint  $c$  is violated, its further mining can be terminated
  - Monotonic: If  $c$  is satisfied, no need to check  $c$  again
  - Succinct<sup>6</sup>: if the constraint  $c$  can be enforced by directly manipulating the data
  - Convertible:  $c$  can be converted to monotonic or anti-monotonic if items can be properly ordered in processing
- Data space pruning constraints
  - Data succinct: Data space can be pruned at the initial pattern mining process
  - Data anti-monotonic: If a transaction  $t$  does not satisfy  $c$ , then  $t$  can be pruned to reduce data processing effort

Anti-monotonic constraints have more pruning power than monotonic constraints.

---

<sup>6</sup>Succinct - краткий

### 5.2.1 Pattern space pruning constraints

Constraint  $c$  is **anti-monotone**: if an itemset  $S$  violates constraint  $c$ , so does any of its superset. That is, mining on itemset  $S$  can be terminated. For example, constraint  $\text{sup}(S) \geq \sigma$  is anti-monotone.

A constraint  $c$  is **monotone**: if an itemset  $S$  satisfies the constraint  $c$ , so does any of its superset. That is, we do not need to check  $c$  in subsequent mining. For example, constraints  $\text{sum}(S.\text{price}) \geq v$  or  $\text{min}(S.\text{price}) \leq v$  are monotone.

### 5.2.2 Data space pruning constraints

A constraint  $c$  is **data anti-monotone**: if a data entry  $t$  cannot satisfy a pattern  $p$  under constraint  $c$ ,  $t$  cannot satisfy  $p$ 's superset either. That's why, data entry  $t$  can be pruned.

**Succinctness**: if the constraint  $c$  can be enforced by directly manipulating the data.

**Convertible constraints**: convert tough<sup>7</sup> constraints into (anti-)monotone by proper ordering of items in transactions. For example, ordering items in value-descending order makes the constraint  $\text{avg}(S.\text{profit}) > 20$  anti-monotone *if the patterns grow in the right order*.

## 5.3 Recommended Readings

- R. Srikant, Q. Vu, and R. Agrawal, «Mining association rules with item constraints», KDD'97
- R. Ng, L.V.S. Lakshmanan, J. Han & A. Pang, Exploratory mining and pruning optimizations of constrained association rules», SIGMOD'98
- G. Grahne, L. Lakshmanan, and X. Wang, «Efficient mining of constrained correlated sets», ICDE'00
- J. Pei, J. Han, and L. V. S. Lakshmanan, «Mining Frequent Itemsets with Convertible Constraints», ICDE'01
- J. Pei, J. Han, and W. Wang, «Mining Sequential Patterns with Constraints in Large Databases», CIKM'02
- F. Bonchi, F. Giannotti, A. Mazzanti, and D. Pedreschi, «ExAnte: Anticipated Data Reduction in Constrained Pattern Mining», PKDD'03
- F. Zhu, X. Yan, J. Han, and P. S. Yu, «gPrune: A Constraint Pushing Framework for Graph Pattern Mining», PAKDD'07

## 6 Sequential Pattern Mining

### 6.1 Sequential Pattern

Sequence  $\rightarrow$  Element  $\rightarrow$  Item or Event (items within an element are unordered)

---

<sup>7</sup>Tough - жесткий

The Apriori property still holds: if a subsequence  $s_1$  is infrequent, none of  $s_1$ 's super-sequences can be frequent.

Algorithms:

- Generalized Sequential Patterns: **GSP**
- Vertical format-based mining: **SPADE**
- Pattern-growth methods: **PrefixSpan**
- Mining closed sequential patterns: **CloSpan**

## 6.2 GSP: Apriori-Based Sequential Pattern Mining

---

### Algorithm 4 GSP

---

```

k = 1
repeat
  find length=k frequent sequences
  Apriori: remove candidates with sup < min_sup
  length=k frequent sequences  $\Rightarrow$  length=(k+1) candidate sequences
  k = k + 1
until no frequent sequences or candidates

```

---

## 6.3 SPADE: Sequential Pattern Mining in Vertical Data Format

**SPADE** = Sequential Pattern Discovery using Equivalent Class

SID	Sequence
1	<a( <u>abc</u> )(a <u>c</u> )d(cf)>
2	<(ad)c(bc)(ae)>
3	<(ef)( <u>ab</u> )(df) <u>cb</u> >
4	<eg(af)cbc>

*min\_sup = 2*

Figure 4: A sequence database

SID	EID	Items
1	1	a
1	2	abc
1	3	ac
1	4	d
1	5	cf
2	1	ad
2	2	c
2	3	bc
2	4	ae
3	1	ef
3	2	ab
3	3	df
3	4	c
3	5	b
4	1	e
4	2	g
4	3	af
4	4	c
4	5	b
4	6	c

a		b		...
SID	EID	SID	EID	...
1	1	1	2	
1	2	2	3	
1	3	3	2	
2	1	3	5	
2	4	4	5	
3	2			
4	3			

ab			ba		
SID	EID (a)	EID(b)	SID	EID (b)	EID(a)
1	1	2	1	2	3
2	1	3	2	3	4
3	2	5			
4	3	5			

aba					...
SID	EID (a)	EID(b)	EID(a)		
1	1	2	3		
2	1	3	4		

Figure 5: SPADE algorithm

## 6.4 PrefixSpan: A Pattern-Growth Approach

**PrefixSpan** = Prefix-projected Sequential pattern mining

SID	Sequence	Prefix	Suffix (Projection)
10	<a(abc)(ac)d(cf)>	<a>	<(abc)(ac)d(cf)>
20	<(ad)c(bc)(ae)>	<aa>	<(_bc)(ac)d(cf)>
30	<(ef)(ab)(df)cb>	<ab>	<(_c)(ac)d(cf)>
40	<eg(af)cbc>		

Figure 6: SPADE algorithm

PrefixSpan Mining: Prefix Projections

- Step 1: Find length-1 sequential patterns: <a>, <b>, etc.
- Step 2: Divide search space and mine each projected DB: <a>-projected DB, <b>-projected DB, etc.

## 6.5 CloSpan: Mining Closed Sequential Patterns

**Definition.** A closed sequential pattern  $\alpha$ : there exists no superpattern  $\beta$  such that  $\beta$  and  $\alpha$  have the same support:



$$CS = \{\alpha \mid \alpha \in FS \text{ and } \nexists \beta \in FS, \text{ such that } \alpha \subseteq \beta \text{ and } sup(\alpha) = sup(\beta)\}$$

CloSpan is based on this property: if  $s \supset s_1$  then  $s$  is closed only if two project DBs have the same size. So redundant search space can be pruned using **Backward Subpattern** and **Backward Superpattern** pruning.




Figure 7: CloSpan pruning algorithm

## 6.6 Constraint-Based Sequential-Pattern Mining

- **Anti-monotonic:** If  $S$  violates  $c$ , the super-sequences of  $S$  also violate  $c$
- **Monotonic:** If  $S$  satisfies  $c$ , the super-sequences of  $S$  also do so
- **Data anti-monotonic:** If a sequence  $s_1$  with respect to  $S$  violates  $c_3$ ,  $s_1$  can be removed
- **Succinct:** Enforce constraint  $c$  by explicitly manipulating data
- **Convertible:** Projection based on the sorted value not in sequence order

### 6.6.1 Timing-Based Constraints

- **Order constraint:** Some items must happen before the other. Anti-monotonic: constraint-violating sub-patterns pruned
- **Min-gap/max-gap constraint:** Confines two elements in a pattern. Succinct: enforced directly during pattern growth
- **Max-span constraint:** Maximum allowed time difference between the 1st and the last elements in the pattern. Succinct: enforced directly when the 1st element is determined
- **Window size constraint:** Time window allows a group of consecutive elements of a data-sequence to be merged and treated as a single element as long as their timestamps are within the user-specified window-size.

## 6.7 Recommended Readings

- M. N. Garofalakis, R. Rastogi, K. Shim: Mining Sequential Patterns with Regular Expression Constraints. IEEE Trans. Knowl. Data Eng. 14(3), 2002

- H. Mannila, H. Toivonen, and A. I. Verkamo, "Discovery of frequent episodes in event sequences", Data Mining and Knowledge Discovery, 1997
- J. Pei, J. Han, B. Mortazavi-Asl, J. Wang, H. Pinto, Q. Chen, U. Dayal, and M.-C. Hsu, "Mining Sequential Patterns by Pattern-Growth: The PrefixSpan Approach", IEEE TKDE, 16(10), 2004
- J. Pei, J. Han, and W. Wang, "Constraint-based sequential pattern mining: the pattern-growth methods", J. Int. Inf. Sys., 28(2), 2007
- R. Srikant and R. Agrawal, "Mining sequential patterns: Generalizations and performance improvements", EDBT'96
- X. Yan, J. Han, and R. Afshar, "CloSpan: Mining Closed Sequential Patterns in Large Datasets", SDM'03
- M. Zaki, "SPADE: An Efficient Algorithm for Mining Frequent Sequences", Machine Learning, 2001

## 7 Lecture 8. Graph Pattern Mining

### 7.1 Frequent (Sub)Graph Patterns

Given a labeled graph dataset  $D = \{G_1, G_2, \dots, G_n\}$ , the supporting graph set of a sub-graph  $g$  is  $D_g = \{G_i \mid g \subseteq G_i, G_i \in D\}$ :

$$\text{support}(g) = |D_g| / |D|$$

A (sub)graph  $g$  is frequent if  $\text{support}(g) \geq \text{min\_sup}$ .




Figure 8: Example: Chemical structures

## 7.2 Apriori-Based Approach

**The Apriori property** (anti-monotonicity): a size- $k$  subgraph is frequent if and only if all of its subgraphs are frequent.

Candidate generation: a candidate size- $(k+1)$  edge/vertex subgraph is generated if its corresponding two  $k$ -edge/vertex subgraphs are frequent:

- AGM - Generating new graphs with one more vertex
- FSG - Generating new graphs with one more edge (more efficient)

Iterative mining process: Candidate-generation  $\rightarrow$  candidate pruning  $\rightarrow$  support counting  $\rightarrow$  candidate elimination.

## 7.3 gSPAN: Graph Pattern Growth

Depth-first growth of subgraphs from  $k$ -edge to  $(k+1)$ -edge, then  $(k+2)$ -edge subgraphs generates many duplicate subgraphs.

**Right-most path extension** in subgraph pattern growth reduces generation of duplicate subgraphs: *take the path from root to the right-most leaf (choose the vertex with the smallest index at each step)*. The Enumeration of graphs using right-most path extension is complete.

DFS Code: flatten a graph into a sequence using depth-first search




Figure 9: gSPAN

## 7.4 Mining Closed Graph Patterns

A frequent graph  $G$  is closed if there exists no supergraph of  $G$  that carries the same support as  $G$ .

**CloseGraph** algorithm: mining closed graph patterns by extending gSpan. Suppose  $G$  and  $G_1$  are frequent, and  $G$  is a subgraph of  $G_1$ . If in any part of the graph in the dataset where  $G$  occurs,  $G_1$  also occurs, then we need not grow  $G$  (except some special, subtle cases), since *none of  $G$ 's children will be closed except those of  $G_1$* .

## 7.5 gIndex: A Graph Indexing Method




Figure 10: Graph query

- use frequent substructures for indexing
- discriminative substructures: reduce index size by removing similar (not discriminative) substructures from the index

**Definition.** Fragment  $x$  is **discriminative** with respect to feature set  $F$  if  $D_x \ll \bigcap_{f \in F \wedge f \subseteq x} D_f$ , where  $D_x$  is the set of graphs containing  $x$ .

**Selection:** Given a set of indexing features  $f_1, f_2, \dots, f_n$ , and a new structure  $x$  ( $x$  should be either redundant or discriminative), the extra indexing power is measured by occurrence probability

$$Pr(x \mid f_1, f_2, \dots, f_n) = \frac{\left| \bigcap_{f \in F \wedge f \subseteq x} D_f \right|}{|D_x|}$$

When  $Pr(x \mid f_1, f_2, \dots, f_n) \ll 1$ ,  $x$  is a discriminative structure and should be included in the index.

## 7.6 SpiderMine: Mining Top-K Large Structural Patterns in a Massive Network

**SpiderMine:** mine top-K largest frequent substructure patterns whose diameter is bounded by  $D_{max}$  with a probability at least  $1 - \epsilon$ . General idea: large patterns are composed of a number of small components («spiders») which will eventually connect together after some rounds of pattern growth.

An  $r$ -spider is a frequent graph pattern  $P$  such that there exists a vertex  $u$  of  $P$ , and all other vertices of  $P$  are within distance  $r$  from  $u$ .

The SpiderMine Algorithm

- Mine the set  $S$  of all the  $r$ -spiders
- Randomly draw  $M$   $r$ -spiders

- Grow these  $M$  r-spiders for  $t = D_{max}/2$  iterations, and merge two patterns whenever possible
- Discard unmerged patterns
- Continue to grow the remaining ones to maximum size
- Return the top- $K$  largest ones in the result

SpiderMine general ideas:

- Small patterns are much less likely to be hit in the random draw
- Even if a small pattern is hit, it is even less likely to be hit multiple times
- The larger the pattern, the greater the chance it is hit and saved

## 7.7 Recommended Readings

- C. Borgelt and M. R. Berthold, «Mining molecular fragments: Finding relevant substructures of molecules», ICDM'02
- J. Huan, W. Wang, and J. Prins. «Efficient mining of frequent subgraph in the presence of isomorphism», ICDM'03
- A. Inokuchi, T. Washio, and H. Motoda. «An apriori-based algorithm for mining frequent substructures from graph data», PKDD'00
- M. Kuramochi and G. Karypis. «Frequent subgraph discovery», ICDM'01
- S. Nijssen and J. Kok. A quickstart in frequent structure mining can make a difference. KDD'04
- N. Vanetik, E. Gudes, and S. E. Shimony. «Computing frequent graph patterns from semistructured data», ICDM'02
- X. Yan and J. Han, «gSpan: Graph-Based Substructure Pattern Mining», ICDM'02
- X. Yan and J. Han, «CloseGraph: Mining Closed Frequent Graph Patterns», KDD'03
- X. Yan, P. S. Yu, and J. Han, «Graph Indexing: A Frequent Structure-based Approach», SIGMOD'04
- F. Zhu, Q. Qu, D. Lo, X. Yan, J. Han, and P. S. Yu, «Mining Top- $K$  Large Structural Patterns in a Massive Network», VLDB'11

## 8 Lecture 9. Pattern-Based Classification

### 8.1 Classification: Basic Concepts

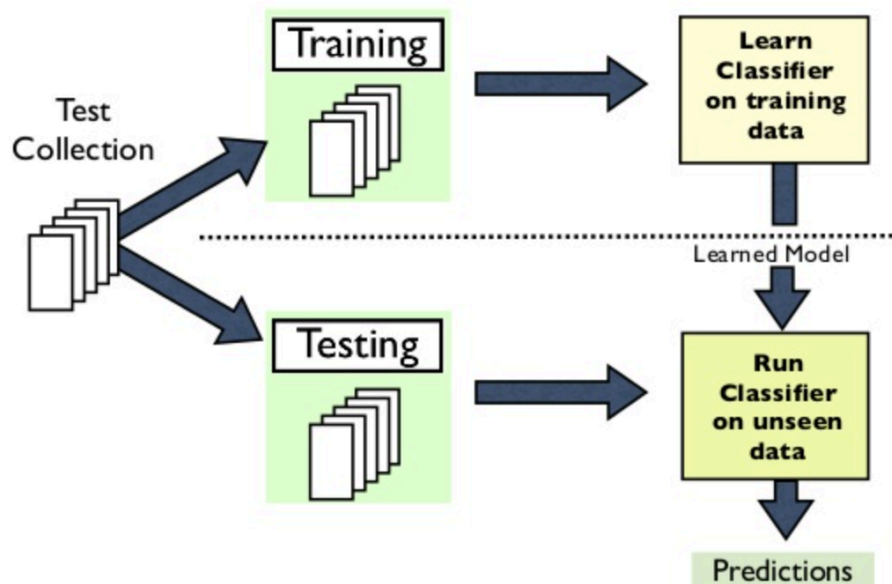


Figure 11: Classification

Typical Classification Methods:

- Support Vector Machines
- Decision Tree
- Neural Network
- Bayesian Network

### 8.2 Pattern-based classification methods

- CBA [Liu, Hsu & Ma, KDD'98]: Use high-conf., high-support class association rules to build classifiers
- Emerging patterns [Dong & Li, KDD'99]: Patterns whose support changes significantly between the two classes
- CMAR [Li, Han & Pei, ICDM'01]: Multiple rules in prediction
- CPAR [Yin & Han, SDM'03]: Beam search on multiple prediction rules
- RCBT [Cong et al., SIGMOD'05]: Build classifier based on mining top-k covering rule groups with row enumeration (for high-dimensional data)
- Lazy classifier [Veloso, Meira & Zaki, ICDM'06]: For a test  $t$ , project training data  $D$  on  $t$ , mine rules from  $D_t$ , predict on the best rule
- Discriminative pattern-based classification [Cheng et al., ICDE'07]

### 8.3 CBA: Classification Based on Associations

- Mine high-confidence, high-support class association rules
- LHS: conjunctions of attribute-value pairs; RHS: class labels  
 $p_1 \wedge p_2 \dots \wedge p_l \rightarrow [A_{class-label} = C]$
- Rank rules in descending order of confidence and support
- Classification: Apply the first rule that matches a test case; otherwise apply the default rule

### 8.4 CMAR: Classification Based on Multiple Association Rules

Rule pruning whenever a rule is inserted into the tree:

- Given two rules,  $R_1$  and  $R_2$ , if the antecedent<sup>8</sup> of  $R_1$  is more general than that of  $R_2$  and  $conf(R_1) \geq conf(R_2)$ , then prune  $R_2$
- Prunes rules for which the rule antecedent and class label are not positively correlated, based on the  $\chi^2$  test of statistical significance

Classification based on generated/pruned rules:

- If only one rule satisfies tuple X, assign the class label of the rule
- If a ruleset S satisfies X
  - Divide S into groups according to class labels
  - Use a weighted  $\chi^2$  measure to find the strongest group of rules, based on the statistical correlation of rules within a group
  - Assign X the class label of the strongest group

### 8.5 Discriminative Pattern-Based Classification

**Principle:** Mining discriminative frequent patterns as high-quality features and then apply any classifier.

Framework (**PatClass**):

- Feature construction by **frequent itemset mining**
- Feature selection (e.g., using **Maximal Marginal Relevance (MMR)**)
  - Select **discriminative features** (i.e., that are relevant but minimally similar to the previously selected ones)
  - Remove redundant or closely correlated features
- Model learning: apply a general classifier, such as SVM or C4.5, to build a classification model

---

<sup>8</sup>Antecedent - предшественник

### 8.5.1 On the Power of Discriminative Patterns

K-itemsets are often more informative than single features (1-itemsets) in classification. Computation on real datasets shows: the discriminative power of k-itemsets (for  $k > 1$  but often  $\leq 10$ ) is higher than that of single features.

Computation on real datasets shows: pattern frequency (but not too frequent) is strongly tied with the discriminative power (information gain). Information gain upper bound monotonically increases with pattern frequency.

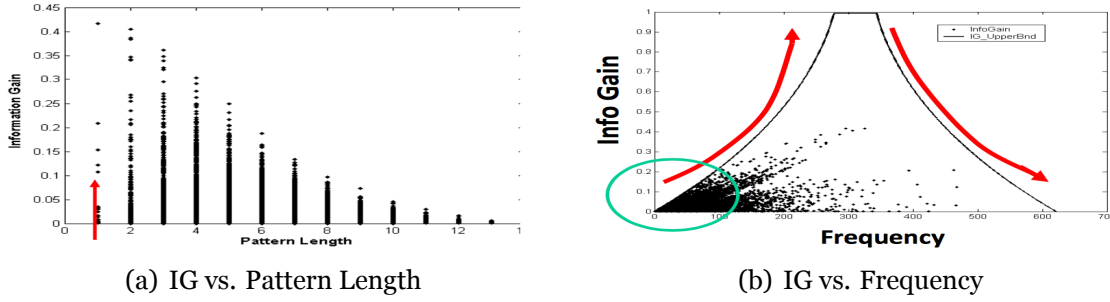


Figure 12: Information Gain

Information gain formula:

$$IG(C | X) = H(C) - H(C | X), \text{ where}$$

$$H(C) = - \sum_{i=1}^m p_i \log_2(p_i) - \text{entropy of given data}$$

$$H(C | X) = \sum_j P(X = x_j) H(Y | X = x_j) - \text{conditional entropy of study focus}$$

## 8.6 DDPMine: Direct Mining of Discriminative Patterns

General methodology:

- Input: A set of training instances D and a set of features F
- Iteratively perform feature selection based on the «**sequential coverage**» paradigm
  - Select the feature  $f_i$  with the highest discriminative power
  - Remove instances  $D_i$  from D covered by the selected feature  $f_i$

Implementation:

- Integration of **branch-and-bound search** with FP-growth mining
- Iteratively eliminate training instances and **progressively shrink the FP-tree**

### 8.6.1 Branch-and-Bound Search

- The discriminative power (information gain) of a low frequency pattern is upper bounded by a small value
- During FPGrowth mining we record the most discriminative itemset discovered so far and its information gain value  $g_{best}$



- Before constructing a conditional FP-tree, we first estimate the upper bound of information gain based on the conditional DB
- If the upper bound value  $\leq g_{best}$ , skip this conditional FP-tree and its subsequent trees

## 8.7 Recommended Readings

- H. Cheng, X. Yan, J. Han & C.-W. Hsu, Discriminative Frequent Pattern Analysis for Effective Classification, ICDE'07
- H. Cheng, X. Yan, J. Han & P. S. Yu, Direct Discriminative Pattern Mining for Effective Classification, ICDE'08
- G. Cong, K. Tan, A. Tung & X. Xu. Mining Top-k Covering Rule Groups for Gene Expression Data, SIGMOD'05
- M. Deshpande, M. Kuramochi, N. Wale & G. Karypis. Frequent Substructure-based Approaches for Classifying Chemical Compounds, TKDE'05
- G. Dong & J. Li. Efficient Mining of Emerging Patterns: Discovering Trends and Differences, KDD'99
- W. Fan, K. Zhang, H. Cheng, J. Gao, X. Yan, J. Han, P. S. Yu & O. Verscheure. Direct Mining of Discriminative and Essential Graphical and Itemset Features via Model-based Search Tree, KDD'08
- W. Li, J. Han & J. Pei. CMAR: Accurate and Efficient Classification based on Multiple Class-association Rules, ICDM'01
- B. Liu, W. Hsu & Y. Ma. Integrating Classification and Association Rule Mining, KDD'98
- J. Wang and G. Karypis. HARMONY: Efficiently Mining the Best Rules for Classification, SDM'05
- X. Yin & J. Han. CPAR: Classification Based on Predictive Association Rules, SDM'03

## 9 Lecture 10. Exploring Pattern Mining Applications

Frequent Pattern Mining for Text Data — Phrase Mining and Topic Modeling:

- Strategy 1: Simultaneously Inferring Phrases and Topics
  - Generate bag-of-words  $\rightarrow$  generate sequence of tokens
  - Bigram topical model [Wallach'06], topical n-gram model [Wang, et al.'07], phrase discovering topic model [Lindsey, et al.'12]
- Strategy 2: Post Topic Modeling Phrase Construction
  - Post bag-of-words model inference, visualize topics with n-grams
  - Label topic [Mei et al.'07], TurboTopic [Blei & Lafferty'09], KERT [Danilevsky, et al.'14]
- Strategy 3: First Phrase Mining then Topic Modeling

- Prior bag-of-words model inference, mine phrases and impose on the bag-of-words model
- ToPMine [El-kishky, et al.'15]

## 9.1 Strategy 1: Simultaneously Inferring Phrases and Topics

- **Bigram Topic Model** [Wallach'06]
  - Probabilistic generative model that conditions on previous word and topic when drawing next word
- **Topical N-Grams (TNG)** [Wang, et al.'07]
  - Probabilistic model that generates words in textual order
  - Create n-grams by concatenating successive bigrams (a generalization of Bigram Topic Model)
- **Phrase-Discovering LDA<sup>9</sup> (PDLDA)** [Lindsey, et al.'12]
  - Viewing each sentence as a time-series of words, PDLDA posits that the generative parameter (topic) changes periodically
  - Each word is drawn based on previous m words (context) and current phrase topic

## 9.2 Strategy 2: Post Topic Modeling Phrase Construction

### 9.2.1 TurboTopics

**TurboTopics** [Blei & Lafferty'09] – Phrase construction as a post-processing step to Latent Dirichlet Allocation

- Perform Latent Dirichlet Allocation on corpus to assign each token a topic label
- Merge adjacent unigrams with the same topic label by a distribution-free permutation test on arbitrary-length back-off model
- End recursive merging when all significant adjacent unigrams have been merged

### 9.2.2 KERT

**KERT** [Danilevsky et al.'14] – Phrase construction as a post-processing step to Latent Dirichlet Allocation

- Perform frequent pattern mining on each topic
- Perform phrase ranking based on four different criteria

Framework of KERT:

- Run bag-of-words model inference and assign topic label to each token
- Extract candidate keyphrases within each topic (frequent pattern mining)
- Rank the keyphrases in each topic

---

<sup>9</sup>LDA - Latent Dirichlet Allocation



## 9.4 Recommended Readings

- M. Danilevsky, C. Wang, N. Desai, X. Ren, J. Guo, J. Han. «Automatic Construction and Ranking of Topical Keyphrases on Collections of Short Documents», SDM'14
- X. Wang, A. McCallum, X. Wei. Topical n-grams: Phrase and topic discovery, with an application to information retrieval, ICDM'07
- R. V. Lindsey, W. P. Headden, III, M. J. Stipicevic. A phrase-discovering topic model using hierarchical pitman-yor processes, EMNLP-CoNLL'12.
- Q. Mei, X. Shen, C. Zhai. Automatic labeling of multinomial topic models, KDD'07
- D. M. Blei and J. D. Lafferty. Visualizing Topics with Multi-Word Expressions, arXiv:0907.1013, 2009
- M. Danilevsky, C. Wang, N. Desai, J. Guo, J. Han. Automatic Construction and Ranking of Topical Keyphrases on Collections of Short Documents, SDM'14
- A. El-Kishky, Y. Song, C. Wang, C. R. Voss, J. Han. Scalable Topical Phrase Mining from Text Corpora, VLDB'15
- K. Church, W. Gale, P. Hanks, D. Hindle. Using Statistics in Lexical Analysis. In U. Zernik (ed.), Lexical Acquisition: Exploiting On-Line Resources to Build a Lexicon. Lawrence Erlbaum, 1991

## 10 Lecture 11: Advanced Topics on Pattern Discovery

### 10.1 Frequent Pattern Mining in Data Streams

#### 10.1.1 Lossy Counting Algorithm




Figure 15: Lossy Counting Algorithm

Given:

- support threshold =  $\sigma$
- error threshold =  $\varepsilon$
- stream length =  $N$

Output: items with frequency counts exceeding  $(\sigma - \varepsilon) \times N$

$$\text{frequency count error} \leq \text{number of buckets} = \frac{N}{\text{bucket size}} = \frac{N}{1/\varepsilon} = \varepsilon N$$

Approximation guarantee:

- No false negatives
- False positives have true frequency count at least  $(\sigma - \varepsilon) \times N$
- Frequency count underestimated by at most  $\varepsilon N$

### 10.1.2 Recommended Readings

- G. Manku and R. Motwani, «Approximate Frequency Counts over Data Streams», VLDB'02
- A. Metwally, D. Agrawal, and A. El Abbadi, «Efficient Computation of Frequent and Top-k Elements in Data Streams», ICDT'05

## 10.2 Spatiotemporal and Trajectory Pattern Mining

### 10.2.1 Mining Relative Movement Patterns

- **Flock:** At least  $m$  entities are within a circular region of radius  $r$  and move in the same direction
- **Convoy:** Uses density-based clustering at each timestamp; no need to be a rigid circle
- **Swarm:** Moving objects may not be close to each other for all the consecutive time stamps

### 10.2.2 Recommended Readings

- Y. Huang, S. Shekhar, H. Xiong, Discovering colocation patterns from spatial data sets: A general approach, IEEE Trans. on Knowledge and Data Engineering, 16(12), 2004
- K. Koperski, J. Han, «Discovery of Spatial Association Rules in Geographic Information Databases», SSD'95
- Z. Li, B. Ding, J. Han, R. Kays, «Swarm: Mining Relaxed Temporal Moving Object Clusters», VLDB'10
- Z. Li, B. Ding, J. Han, Roland Kays, Peter Nye, «Mining Periodic Behaviors for Moving Objects», KDD'10
- C. Zhang, J. Han, L. Shou, J. Lu, T. La Porta, «Splitter: Mining Fine-Grained Sequential Patterns in Semantic Trajectories», VLDB'14
- Y. Zheng and X. Zhou, Computing with Spatial Trajectories, Springer, 2011

## 10.3 Pattern Discovery for Software Bug Mining

### 10.3.1 Typical Software Bug Detection Methods

- Mining rules from source code
  - Bugs as deviant behavior (e.g., by statistical analysis)
  - Mining programming rules (e.g., by frequent itemset mining)
  - Mining function precedence protocols (e.g., by frequent subsequence mining)
  - Revealing neglected conditions (e.g., by frequent itemset/subgraph mining)
- Mining rules from revision histories
  - By frequent itemset mining
- Mining copy-paste patterns from source code
  - Find copy-paste bugs (e.g., CP-Miner [Li et al., OSDI'04])
  - Reference: Z. Li, S. Lu, S. Myagmar, Y. Zhou, «CP-Miner: A Tool for Finding Copy-paste and Related Bugs in Operating System Code», OSDI'04

### 10.3.2 Mining Copy-and-Paste Bugs

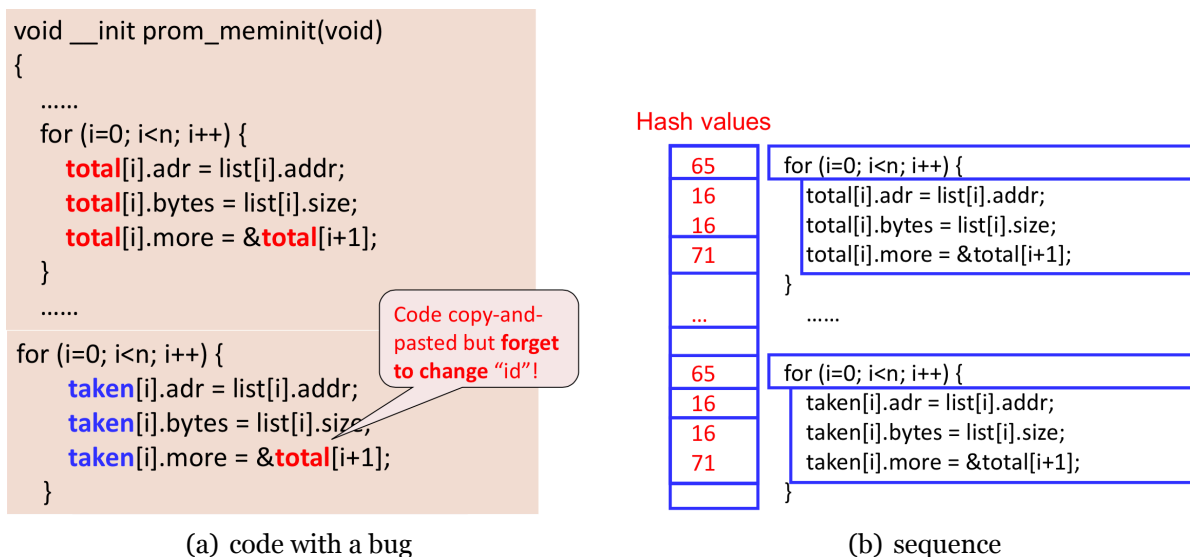


Figure 16: Copy-and-Paste Bugs

- Map each statement to number
- Tokenize each component
  - Different operators, constants, keywords = different tokens
  - Same type of identifiers = same token
- Program = long sequence
  - Cut the long sequence by blocks.

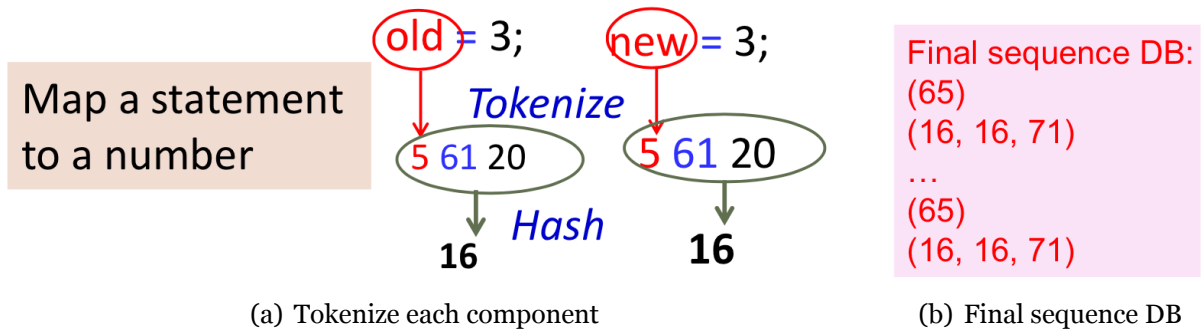


Figure 17: Building Sequence Database from Source Code




Figure 18: Detecting «Forget-to-Change» Bugs

- Modification to the sequence pattern mining algorithm
  - Constrain the max gap
- Composing Larger Copy-Pasted Segments
  - Combine the neighboring copy-pasted segments repeatedly
- Find conflicts: Identify names that cannot be mapped to the corresponding ones
  - If  $0 < \text{unchanged ratio} < \text{threshold}$ , then report it as a bug

## 10.4 Pattern Discovery for Image Analysis

### 10.4.1 Recommended Readings

- Hongxing Wang, Gangqiang Zhao, Junsong Yuan, Visual pattern discovery in image and video data: a brief survey, Wiley Interdisciplinary Review: Data Mining and Knowledge Discovery 4(1): 24-37 (2014)
- Hongxing Wang, Junsong Yuan, Ying Wu, Context-Aware Discovery of Visual Co-Occurrence Patterns. IEEE Transactions on Image Processing 23(4): 1805-1819 (2014)
- Gangqiang Zhao, Junsong Yuan, Discovering Thematic Patterns in Videos via Cohesive Sub-graph Mining. ICDM 2011: 1260-1265
- Junsong Yuan, Ying Wu, Ming Yang, From frequent itemsets to semantically meaningful visual patterns. KDD 2007: 864-873